

# 100 Days with TensorFlow - One Day Per Day!

This document summarizes the 100 lessons on learning TensorFlow, one lesson each day. The difficulty will gradually increase, with weekly progress reviews. The goal is to build a solid foundation in TensorFlow and work on a final project together. Each lesson includes code examples and explanations.

## Lesson 1: Introduction to TensorFlow Basics

In this lesson, we introduced TensorFlow and its use of tensors for handling data. Tensors are the fundamental building blocks in TensorFlow. A tensor is a generalization of vectors and matrices to potentially higher dimensions. TensorFlow represents all data in this multi-dimensional array format.

We started by printing the TensorFlow version and creating two constant tensors, using `'tf.constant()'`. Tensors in TensorFlow are immutable (their values can't change), and they can store data of any type like integers, floats, strings, or booleans.

We then performed a basic addition operation using `'tf.add()'`. TensorFlow provides built-in functions for performing mathematical operations on tensors, such as addition, multiplication, and matrix operations. The result of an operation is returned as another tensor.

In Part 2, we explored tensor shapes. A tensor's shape defines its dimensionality, or how many elements it contains in each dimension. A scalar (0-D tensor) has no dimensions, while a vector (1-D tensor) has one. A matrix (2-D tensor) has two dimensions, and higher-dimensional tensors can represent more complex data structures like batches of images.

Reshaping tensors is an important concept in TensorFlow. It allows you to change the shape of the

tensor without changing its underlying data. For example, a tensor with 4 elements can be reshaped into a 2x2 matrix or a 1D vector. However, the total number of elements must remain the same.

In Part 3, we applied various tensor operations. TensorFlow supports arithmetic operations like addition, subtraction, multiplication, and division directly on tensors. We also performed matrix multiplication using 'tf.matmul()', which is useful in neural networks, where weights and inputs are often represented as matrices.

Finally, we discussed why reshaping is important. Reshaping tensors helps adjust the data's structure for specific machine learning models, batch processing, or to meet the input/output requirements of layers in neural networks.

### **Fun Facts**

**Did you know? When reshaping tensors, the total number of elements must remain the same! For example, a tensor with 4 elements can be reshaped into (2, 1, 2), but not (2, 1, 1), since the number of elements would not match. Tensors can be reshaped to fit various dimensions based on how you need to structure the data.**

### **Code Used in Lesson 1**

#### **# Lesson 1, Part 1 - Basic TensorFlow Program**

```
import tensorflow as tf

print("TensorFlow version:", tf.__version__)

tensor_a = tf.constant(5)

tensor_b = tf.constant(3)

result = tf.add(tensor_a, tensor_b)
```

```
print("Result of addition:", result)
```

### **# Lesson 1, Part 2 - Tensor Types and Shapes**

```
int_tensor = tf.constant(10, dtype=tf.int32)
```

```
float_tensor = tf.constant(10.5, dtype=tf.float32)
```

```
string_tensor = tf.constant("Hello TensorFlow")
```

```
bool_tensor = tf.constant(True)
```

```
print("Integer Tensor:", int_tensor)
```

```
print("Float Tensor:", float_tensor)
```

```
print("String Tensor:", string_tensor)
```

```
print("Boolean Tensor:", bool_tensor)
```

### **# Reshaping a tensor**

```
reshaped_tensor = tf.reshape(tensor_b, (2, 1))
```

```
print("Reshaped Tensor:", reshaped_tensor)
```

### **# Lesson 1, Part 3 - Tensor Operations**

```
add_result = tf.add(tensor_a, tensor_b)
```

```
multiply_result = tf.multiply(tensor_a, tensor_b)
```

```
subtract_result = tf.subtract(tensor_a, tensor_b)
```

```
divide_result = tf.divide(tensor_a, tensor_b)
```

```
print("Addition Result:", add_result)
```

```
print("Multiplication Result:", multiply_result)
```

```
print("Subtraction Result:", subtract_result)
```

```
print("Division Result:", divide_result)
```

### **# Matrix multiplication**

```
matrix_1 = tf.constant([[1, 2], [3, 4]])
```

```
matrix_2 = tf.constant([[5, 6], [7, 8]])
```

```
matrix_mult_result = tf.matmul(matrix_1, matrix_2)
```

```
print("Matrix Multiplication Result:", matrix_mult_result)
```

### **# Other operations**

```
sum_result = tf.reduce_sum(matrix_1)
```

```
mean_result = tf.reduce_mean(matrix_1)
```

```
max_index = tf.argmax(tensor_b)
```

```
print("Sum of elements:", sum_result)
```

```
print("Mean of elements:", mean_result)
```

```
print("Index of max value:", max_index)
```

## Lesson 2: TensorFlow Operations and Variables

In this lesson, we explored TensorFlow variables, operations, and broadcasting. We started by learning how to create variables in TensorFlow using `'tf.Variable()'`, which allows for mutable values, unlike constant tensors. Variables can be updated with methods like `'.assign()'`.

Next, we delved into basic operations beyond addition, such as matrix multiplication (`'tf.matmul()'`) and element-wise multiplication (`'tf.multiply()'`), to perform matrix operations commonly needed in machine learning and neural networks.

Finally, we learned about broadcasting, a powerful TensorFlow feature that allows tensors of different shapes to be used together in element-wise operations. Broadcasting stretches the smaller tensor to match the shape of the larger tensor, allowing for efficient operations.

### Fun Facts

**Matrix multiplication and element-wise multiplication serve different purposes. Matrix multiplication combines rows and columns and is used in applications such as neural networks and transformations. Element-wise multiplication operates directly on individual elements, useful for scaling or masking. Both are key operations in machine learning.**

### Matrix Multiplication vs. Element-wise Multiplication

**Matrix multiplication (dot product) involves combining rows of one matrix with columns of another. It is widely used in neural networks.**

**Element-wise multiplication multiplies corresponding elements of two matrices directly. Both require specific rules:**

### Matrix Multiplication:

- Requires the number of columns in the first matrix to equal the number of rows in the second.
- Produces a new matrix whose shape is determined by the outer dimensions of the input matrices.

#### **Element-wise Multiplication:**

- Requires matrices to have the same shape.
- Multiplies corresponding elements of the two matrices directly.

#### **Code Examples for Practice**

##### **# 1. Matrix Multiplication**

```
import tensorflow as tf

matrix_a = tf.constant([[1, 2], [3, 4]])
matrix_b = tf.constant([[5, 6], [7, 8]])
result = tf.matmul(matrix_a, matrix_b)
print(result.numpy())
```

##### **# 2. Element-wise Multiplication**

```
elementwise_product = tf.multiply(matrix_a, matrix_b)
print(elementwise_product.numpy())
```

##### **# 3. Broadcasting Example**

```
tensor_a = tf.constant([1, 2])
tensor_b = tf.constant([[1, 2], [3, 4]])
broadcast_result = tf.add(tensor_a, tensor_b)
```

```
print(broadcast_result.numpy())
```

## Fun Facts

### 1. Broadcasting Tensors Is Like Sharing Hats:

Imagine you have a group of people, and they all need hats. Instead of making a custom hat for each person, you use a single hat size and broadcast it to everyone. In TensorFlow, broadcasting works similarly-smaller tensors are stretched across larger ones, allowing element-wise operations without reshaping.

### 2. Matrix Multiplication vs. Element-wise Multiplication:

Matrix multiplication and element-wise multiplication serve different purposes. Matrix multiplication combines rows and columns, while element-wise multiplication simply multiplies corresponding elements directly.

### 3. TensorFlow's Love for NumPy:

Even though TensorFlow is its own powerful system, it loves working with NumPy. That's why you often see `'.numpy()'` in TensorFlow code. It converts tensors into familiar NumPy arrays, making it easy to switch between both libraries.

### 4. Reshaping Tensors Is Like Rearranging Lego Blocks:

Imagine you're building with Lego blocks. You can rearrange them into different shapes, but the total number of blocks stays the same. Similarly, reshaping tensors keeps the number of elements constant while changing their layout.

### 5. The Magic of Element-wise Broadcasting:

Broadcasting can turn a small tensor into a powerful tool. For example, a 1D tensor like `[1, 2]` can be broadcast across a 2D matrix like `[[1, 2], [3, 4]]`, allowing operations to apply across rows efficiently.

## Lesson 3: TensorFlow Datasets and Data Pipelines

In this lesson, we dive into TensorFlow's `tf.data` API, which is essential for building efficient data pipelines. This API allows us to load, preprocess, and feed data to our models effectively, especially when working with large datasets.

### Part 1: TensorFlow Datasets

The `tf.data.Dataset` class allows you to create datasets from various data sources like lists, tensors, or files. Here's how to create a dataset:

```
import tensorflow as tf

# Create a simple dataset from a list
dataset = tf.data.Dataset.from_tensor_slices([1, 2, 3, 4, 5])

for element in dataset:
    print(element.numpy())
```

### Part 2: Batching and Shuffling Data

Batching groups data into smaller units for efficient processing. Shuffling ensures randomness in the data order to prevent overfitting. Here's how:

```
# Batch the dataset

batched_dataset = dataset.batch(2)

for batch in batched_dataset:
```



```
print(batch.numpy())
```

```
# Shuffle and batch the dataset
```

```
shuffled_dataset = dataset.shuffle(buffer_size=5).batch(2)
```

```
for batch in shuffled_dataset:
```

```
    print(batch.numpy())
```

### **Part 3: Data Preprocessing with Pipelines**

You can apply transformations to datasets with the `map()` function. Here's an example:

```
# Define a simple preprocessing function
```

```
def preprocess(x):
```

```
    return x * 2
```

```
# Apply the preprocessing function
```

```
processed_dataset = dataset.map(preprocess)
```

```
for element in processed_dataset:
```

```
    print(element.numpy())
```

### **Fun Fact: 20 Key TensorFlow Submodules**

1. `tf.data.Dataset` - For creating efficient data pipelines.
2. `tf.keras` - TensorFlow's high-level API for neural networks.
3. `tf.keras.layers.Dense` - A densely-connected neural network layer.
4. `tf.Variable` - Mutable tensors used for storing model parameters.

5. `tf.Tensor` - Represents multi-dimensional arrays.
6. `tf.nn.relu` - Rectified Linear Unit activation function.
7. `tf.image` - Image processing utilities.
8. `tf.linalg.matmul` - Matrix multiplication function.
9. `tf.random` - For generating random numbers.
10. `tf.keras.optimizers.Adam` - Optimizer for training models.
11. `tf.losses.MeanSquaredError` - Loss function for regression tasks.
12. `tf.metrics.Accuracy` - For measuring model accuracy.
13. `tf.summary` - For logging info to TensorBoard.
14. `tf.function` - Converts Python functions into TensorFlow graphs.
15. `tf.saved_model` - For saving and loading models.
16. `tf.keras.models.Sequential` - A simple way to build neural networks.
17. `tf.train.Checkpoint` - For saving and restoring checkpoints.
18. `tf.cast` - Converts tensors to different data types.
19. `tf.distribute.MirroredStrategy` - For distributed training.
20. `tf.ragged.RaggedTensor` - For handling tensors with varying lengths.

## FAQ

Q1: Why is the function name `tf.data.Dataset.from_tensor_slices` so long?

A1: TensorFlow uses hierarchical naming to make it clear which part of the library you are working with. In this case, `'tf.data'` refers to the data module, `'Dataset'` is the class, and `'from_tensor_slices'` is the method that creates a dataset from tensors.

Q2: What happens if I set `buffer_size` to 2?

A2: If `buffer_size` is 2, TensorFlow will shuffle the data in pairs, limiting the randomness to small

groups of elements. The last element may remain unshuffled if it doesn't have a pair.

Q3: Can I shuffle the dataset multiple times in the same buffer?

A3: No, you can't shuffle the dataset multiple times within the same buffer directly. However, you can chain the `shuffle()` method or use the `repeat()` function to achieve a similar effect.

Q4: What is the difference between matrix multiplication and element-wise multiplication?

A4: Matrix multiplication combines rows and columns of two matrices, while element-wise multiplication directly multiplies corresponding elements. Both are important in different neural network operations.

# Lesson 4: Variables, Gradients, and Auto Differentiation in TensorFlow

In this lesson, we learn about how TensorFlow handles variables, computes gradients, and automatically performs differentiation. Gradients are essential in training machine learning models, especially when minimizing loss functions.

## Part 1: Variables in TensorFlow

In TensorFlow, variables are mutable, which means their values can change during training. This is important because model parameters (weights and biases) need to be updated.

```
import tensorflow as tf

# Create a variable
my_variable = tf.Variable(10)

# Modify the variable
my_variable.assign(15)
my_variable.assign_add(5)
```

## Part 2: Gradients and Auto Differentiation

Gradients are crucial for updating model parameters during training. Using TensorFlow's GradientTape, we can compute the derivative of a function.

```
import tensorflow as tf
```

```
def my_function(x):  
    return x ** 2  
  
x = tf.Variable(3.0)  
  
# Compute the gradient of the function with respect to x  
with tf.GradientTape() as tape:  
    y = my_function(x)  
  
dy_dx = tape.gradient(y, x)  
  
print("Gradient of y with respect to x:", dy_dx.numpy())
```

### Part 3: Gradient Descent Example

In this example, we use gradient descent to minimize the function  $f(x) = (x - 7)^2$ . This process uses the gradients to iteratively update  $x$  and minimize the loss.

```
def my_function(x):  
    return (x - 7) ** 2  
  
learning_rate = 0.05  
  
x = tf.Variable(10.0)  
  
for i in range(30):  
    with tf.GradientTape() as tape:
```

```
function = my_function(x)
```

```
grad = tape.gradient(function, x)
```

```
x.assign_sub(learning_rate * grad)
```

```
print(f"Iteration {i+1}: x = {x.numpy()}, function = {function.numpy()}")
```

## Challenges

Challenge 1: Create a variable 'z' and subtract 3 from it repeatedly until its value is less than 5.

Solution:

```
z = tf.Variable(15.0)
```

```
for i in range(20):
```

```
    z.assign_sub(3)
```

```
    if z < 5:
```

```
        break
```

```
print(f"Iteration {i+1}: z = {z.numpy()}")
```

Challenge 2: Modify the function to return  $x^3 + 2x^2 + 5x$  and compute the gradient at  $x=2$ .

Solution:

```
def my_function(x):
```

```
    return x ** 3 + 2 * (x ** 2) + 5 * x
```

```
x = tf.Variable(2.0)
```

```
with tf.GradientTape() as tape:
```

```
    y = my_function(x)
```

```
dy_dx = tape.gradient(y, x)
```

```
print(f"Gradient at x = 2: {dy_dx.numpy()}")
```

## Fun Fact: Automatic Differentiation

TensorFlow's automatic differentiation uses reverse-mode differentiation, which is highly efficient when computing gradients for machine learning models. It helps minimize loss functions quickly.

## Real-life Examples of Minimizing Loss Functions

1. **Image Recognition**: A model predicts what's in an image (e.g., dog or cat). The loss function measures the error between the prediction and the true label. Minimizing the loss helps the model make more accurate predictions.
2. **Stock Price Prediction**: Predicting stock prices requires minimizing the difference between predicted prices and actual prices. The model's parameters are adjusted using gradients.
3. **Speech Recognition**: The model predicts spoken words, and minimizing the loss function reduces the error between the predicted words and the actual transcript.

## FAQ

Q1: What's the difference between a TensorFlow tensor and a variable?

A TensorFlow tensor is immutable, while a variable is mutable, meaning it can be updated during training.

Q2: How does TensorFlow compute gradients automatically?

Using a feature called automatic differentiation, TensorFlow's GradientTape tracks operations and computes gradients.

Q3: Why do we aim to minimize the loss function?

Minimizing the loss function makes the model's predictions more accurate by reducing the error between the predicted and true values.



## Lesson 5: Linear Regression with TensorFlow

In this lesson, we explore linear regression, a fundamental machine learning model. We will implement a simple linear regression model in TensorFlow, understand how gradient descent works, and use the trained model to make predictions.

### Part 1: Understanding Linear Regression

Linear regression models the relationship between a dependent variable ( $y$ ) and an independent variable ( $x$ ). The goal is to find the best-fit line:  $y = wx + b$ , where  $w$  is the weight (slope) and  $b$  is the bias (intercept).

### Part 2: Implementing Linear Regression in TensorFlow

We will implement linear regression using TensorFlow. Below is an example of training a simple linear model using gradient descent.

```
import tensorflow as tf

# Generate some synthetic data
X = tf.constant([1, 2, 3, 4, 5], dtype=tf.float32)
Y = tf.constant([2, 4, 6, 8, 10], dtype=tf.float32)

# Initialize weights and bias
w = tf.Variable(0.0)
b = tf.Variable(0.0)
```

```
# Define the linear regression model
```

```
def linear_regression(x):
```

```
    return w * x + b
```

```
# Define the loss function (MSE)
```

```
def loss_fn(y_true, y_pred):
```

```
    return tf.reduce_mean(tf.square(y_true - y_pred))
```

```
# Training loop
```

```
learning_rate = 0.01
```

```
epochs = 100
```

```
for epoch in range(epochs):
```

```
    with tf.GradientTape() as tape:
```

```
        y_pred = linear_regression(X)
```

```
        loss = loss_fn(Y, y_pred)
```

```
    gradients = tape.gradient(loss, [w, b])
```

```
    w.assign_sub(learning_rate * gradients[0])
```

```
    b.assign_sub(learning_rate * gradients[1])
```

```
if epoch % 10 == 0:
```

```
    print(f"Epoch {epoch}: Loss = {loss.numpy()}, w = {w.numpy()}, b = {b.numpy()}")
```

### Part 3: Making Predictions with the Trained Model

After training the model, we can use it to make predictions. Simply provide new input data, and the

model will predict the output based on the learned parameters.

```
# Predicting a value for a new input

new_X = tf.constant([6], dtype=tf.float32)

predicted_y = linear_regression(new_X)

print(f"Predicted value for x=6: {predicted_y.numpy()}")
```

## Challenges

Challenge 1: Modify the synthetic data so that  $Y = 3x + 2$ . Train the model again and observe the new values for  $w$  and  $b$ .

Challenge 2: Implement a new loss function using Mean Absolute Error (MAE) and compare it to MSE.

## Fun Fact: Gradient Descent

Gradient descent was introduced by mathematician Augustin-Louis Cauchy in the 19th century. It remains one of the most important optimization algorithms in machine learning, powering everything from neural networks to logistic regression models.

## Quiz: TensorFlow Fundamentals (Lessons 1-5)

1. What is a tensor in TensorFlow?

- A) A single-dimensional number
- B) A multi-dimensional array
- C) A type of neural network
- D) A constant value that cannot change

2. Which function is used to calculate gradients in TensorFlow?

- A) `tf.Variable()`
- B) `tf.GradientTape()`
- C) `tf.math.reduce_sum()`
- D) `tf.keras.Model()`

3. What is the purpose of a loss function in machine learning?

- A) To make predictions
- B) To track progress during training
- C) To calculate the difference between true values and predicted values
- D) To initialize model parameters

4. In the linear regression model  $y = wx + b$ , what does  $w$  represent?

- A) Bias
- B) Weight (slope)
- C) Predicted output
- D) Loss

5. Which of the following is an example of a supervised learning task?

- A) Clustering similar data points
- B) Predicting house prices based on features
- C) Reducing the number of features in a dataset
- D) Detecting anomalies in a dataset

6. What is the difference between Mean Squared Error (MSE) and Mean Absolute Error (MAE)?

- A) MSE is more sensitive to large errors than MAE
- B) MAE is more sensitive to large errors than MSE
- C) MSE and MAE both ignore large errors
- D) MAE is faster to compute than MSE

7. What happens when we use a learning rate that is too high during training?

- A) The model learns very slowly
- B) The model converges to the correct values quickly
- C) The model overshoots the optimal solution and oscillates
- D) The model doesn't learn anything at all

8. What is the main role of gradient descent in training a model?

- A) To initialize model parameters
- B) To reduce the loss function by updating the model's weights and biases
- C) To predict the target variable based on input data
- D) To compute the accuracy of the model

9. Which method in TensorFlow creates a dataset from a list of values?

- A) `tf.data.Dataset.batch()`
- B) `tf.data.Dataset.from_tensor_slices()`
- C) `tf.data.Dataset.shuffle()`
- D) `tf.data.Dataset.map()`

10. What is the main purpose of using the `assign_sub()` function in TensorFlow?

- A) To update a variable by subtracting a value from it
- B) To add a value to a variable

C) To reset the value of a variable

D) To delete a variable

## FAQ

Q1: What is `y_pred`, and how is it different from `y_true`?

A: `y_pred` is the predicted output from the model, while `y_true` is the actual, true value from the dataset.

Q2: Can we use `gradients[1]` for `w` and `gradients[0]` for `b`?

A: No, each index corresponds to a specific variable. `gradients[0]` is for `w`, and `gradients[1]` is for `b`. Swapping them would lead to incorrect updates.

Q3: Why do we aim to minimize the loss function?

A: Minimizing the loss function reduces the error between the model's predictions and the true values, improving model accuracy.

Q4: What happens if we don't know `Y` in supervised learning tasks?

A: If `Y` is unknown, it's an unsupervised learning problem. We use techniques like clustering or dimensionality reduction to find patterns in the data without needing labeled outputs.

Q5: Why do we use learning rates in gradient descent?

A: The learning rate controls the size of the steps we take during gradient descent. A high learning rate can cause overshooting, while a low rate may result in slow convergence.